

KNN AND NAVIE BAIYES

2148059

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Import the Dataset

```
In [3]: #READING DATASET
df=pd.read_csv("/Users/persie/Downloads/diabetes.csv")
```

```
In [5]: df.head()
```

```
Out[5]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFun |
|---|-------------|---------|---------------|---------------|---------|------|---------------------|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 | |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 | |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | |
| 3 | 1 | 89 | 66 | 23 | 94 | 28.1 | |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | |

Exploratory Data Analysis

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              768 non-null    int64
2   BloodPressure                        768 non-null    int64
3   SkinThickness                        768 non-null    int64
4   Insulin                              768 non-null    int64
5   BMI                                  768 non-null    float64
6   DiabetesPedigreeFunction              768 non-null    float64
7   Age                                  768 non-null    int64
8   Outcome                              768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [7]: df.describe()
```

```
Out [7]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | Diabetes |
|-------|-------------|------------|---------------|---------------|------------|------------|------------|
| count | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 |
| mean | 3.845052 | 120.894531 | 69.105469 | 20.536458 | 79.799479 | 31.992578 | 0.349130 |
| std | 3.369578 | 31.972618 | 19.355807 | 15.952218 | 115.244002 | 7.884160 | 0.471405 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 1.000000 | 99.000000 | 62.000000 | 0.000000 | 0.000000 | 27.300000 | 0.000000 |
| 50% | 3.000000 | 117.000000 | 72.000000 | 23.000000 | 30.500000 | 32.000000 | 0.000000 |
| 75% | 6.000000 | 140.250000 | 80.000000 | 32.000000 | 127.250000 | 36.600000 | 0.000000 |
| max | 17.000000 | 199.000000 | 122.000000 | 99.000000 | 846.000000 | 67.100000 | 1.000000 |

```
In [8]: df.isna().sum()
```

```
Out [8]: Pregnancies      0
Glucose      0
BloodPressure  0
SkinThickness 0
Insulin      0
BMI          0
DiabetesPedigreeFunction 0
Age          0
Outcome      0
dtype: int64
```


No Null Values

```
In [9]: df['Outcome'].value_counts()
```

```
Out [9]: 0    500
1     268
Name: Outcome, dtype: int64
```

```
In [10]: sns.countplot(x = df['Outcome'], data = df)
plt.xticks(rotation = 90)
```

```
Out[10]: (array([0, 1]), [Text(0, 0, '0'), Text(1, 0, '1')])
```



The target is imbalanced

```
In [11]: # separate x and y from the dataset
X = df.drop("Outcome", axis = 1)
y = df.Outcome.values
```

```
In [12]: from imblearn.over_sampling import SMOTE
sm = SMOTE(random_state=2)
X_new, y_new = sm.fit_resample(X, y.ravel())
```

```
In [13]: print(len(y_new))
print(len(X_new))
```

```
1000
1000
```

```
In [14]: unique, counts = np.unique(y_new, return_counts=True)
dict(zip(unique, counts))
```

```
Out[14]: {0: 500, 1: 500}
```

The class is balanced now

StandardScaler

```
In [15]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
# transform data
X_sc = scaler.fit_transform(X_new)
print(X_sc)
```

```
[[ 0.60196414  0.70922357  0.14749408 ...  0.1281228  0.41056721
  1.43824553]
 [-0.89248506 -1.23707102 -0.15925271 ... -0.81868657 -0.41664108
 -0.2520288 ]
 [ 1.19974381  1.79049834 -0.26150164 ... -1.26503956  0.54543813
 -0.16306699]
 ...
 [ 1.19974381  0.74011713  0.65873874 ...  0.51699547  1.44720258
  1.0823983 ]
 [-0.29470538 -0.31026407  0.50536534 ... -0.21763017 -0.23789621
  2.06097817]
 [-1.1913749  -0.86634824  0.55648981 ...  0.26228036 -0.47817362
 -0.96372325]]
```

Train Test Split


```
In [17]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_sc, y_new, te
```

```
In [18]: sns.countplot(y_train, data = df)
plt.xticks(rotation = 90)
```

/Users/persie/opt/anaconda3/lib/python3.8/site-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[18]: (array([0, 1]), [Text(0, 0, '0'), Text(1, 0, '1')])
```



KNN

```
In [19]: #Fitting K-NN classifier to the training set  
from sklearn.neighbors import KNeighborsClassifier  
clf_knn = KNeighborsClassifier(n_neighbors=5, metric='minkowski', p  
clf_knn.fit(X_train, y_train)
```

```
Out[19]: KNeighborsClassifier()
```

```
In [21]: #Predicting the test set result  
y_pred = clf_knn.predict(X_test)
```

```
In [22]: #Creating the Confusion matrix  
from sklearn.metrics import confusion_matrix, accuracy_score  
cm = confusion_matrix(y_test, y_pred)  
print(cm)  
  
print('Accuracy -> ' + str(accuracy_score(y_test, y_pred)))
```

```
[[82 23]  
 [10 85]]  
Accuracy -> 0.835
```

```

In [23]: from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

# generate a no skill prediction (majority class)
ns_probs = [0 for _ in range(len(y_test))]

# predict probabilities
lr_probs = clf_knn.predict_proba(X_test)
# keep probabilities for the positive outcome only
lr_probs = lr_probs[:, 1]

# calculate scores
ns_auc = roc_auc_score(y_test, ns_probs)
lr_auc = roc_auc_score(y_test, lr_probs)

# summarize scores
print('No Skill: ROC AUC=%.3f' % (ns_auc))
print('KNN: ROC AUC=%.3f' % (lr_auc))

# calculate roc curves
ns_fpr, ns_tpr, _ = roc_curve(y_test, ns_probs)
lr_fpr, lr_tpr, _ = roc_curve(y_test, lr_probs)

# plot the roc curve for the model
plt.plot(ns_fpr, ns_tpr, linestyle='--', label='No Skill')
plt.plot(lr_fpr, lr_tpr, marker='.', label='KNN')


# axis labels
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')

# show the legend
plt.legend()
# show the plot
plt.show()

```

No Skill: ROC AUC=0.500

KNN: ROC AUC=0.901



NAIVE BAYES

```
In [24]: from sklearn.naive_bayes import GaussianNB
         clf_nb = GaussianNB()
         clf_nb.fit(X_train,y_train)
         ypred = clf_nb.predict(X_test)
```

```
In [25]: from sklearn.metrics import confusion_matrix, accuracy_score, class
         cm = confusion_matrix(y_pred,y_test)
         print(cm)

[[82 10]
 [23 85]]
```

```
In [26]: acc = accuracy_score(y_pred,y_test)
         print(acc)

0.835
```

```
In [27]: cr = classification_report(y_pred,y_test)
         print(cr)
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.78 | 0.89 | 0.83 | 92 |
| 1 | 0.89 | 0.79 | 0.84 | 108 |
| accuracy | | | 0.83 | 200 |
| macro avg | 0.84 | 0.84 | 0.83 | 200 |
| weighted avg | 0.84 | 0.83 | 0.84 | 200 |

```

In [28]: from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

# generate a no skill prediction (majority class)
ns_probs = [0 for _ in range(len(y_test))]

# predict probabilities
lr_probs = clf_nb.predict_proba(X_test)
# keep probabilities for the positive outcome only
lr_probs = lr_probs[:, 1]

# calculate scores
ns_auc = roc_auc_score(y_test, ns_probs)
lr_auc = roc_auc_score(y_test, lr_probs)

# summarize scores
print('No Skill: ROC AUC=%.3f' % (ns_auc))
print('KNN: ROC AUC=%.3f' % (lr_auc))

# calculate roc curves
ns_fpr, ns_tpr, _ = roc_curve(y_test, ns_probs)
lr_fpr, lr_tpr, _ = roc_curve(y_test, lr_probs)

# plot the roc curve for the model
plt.plot(ns_fpr, ns_tpr, linestyle='--', label='No Skill')
plt.plot(lr_fpr, lr_tpr, marker='.', label='Naive Bayes')

# axis labels
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')

# show the legend
plt.legend()
# show the plot
plt.show()

```

No Skill: ROC AUC=0.500
KNN: ROC AUC=0.854

