

## implementation of regression model

2148059

**1)catherine is very curious about analysing the position of women in this modern era. she has an intuition that many bias exist when it comes to women's economic status. find out whether our assumption is correct or wrong.**

in order to understand the existence of bias, take the help of adult dataset from uci repository, and come up with your conclusion.

*the 2022 women's day theme is breaking the bias, through your research find out at what different the bias exist and what should be the solution to overcome it.*

**2)download any stock data, perform both single and multi linear regression**

```
In [262]: import pandas as pd
import numpy as np

import warnings
warnings.filterwarnings('ignore')
```

```
In [263]: data=pd.read_csv('/Users/persie/Downloads/adult.csv')
```

```
In [264]: data.columns=['age',
'workclass',
'fnlwgt',
'education',
'education-num',
'marital-status',
'occupation',
'relationship',
'race',
'sex',
'capital-gain',
'capital-loss',
'hours-per-week',
'native-country','income' ]
```

```
In [265]: data.head()
```

Out[265]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	income
0	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=5
1	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=5
2	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=5
3	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=5
4	37	Private	284582	Masters	14	Married-civ-spouse	Exec-managerial	Wife	White	Female	0	0	40	United-States	<=5

In [266]: data.describe()

Out[266]:

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week
<b>count</b>	32560.000000	3.256000e+04	32560.000000	32560.000000	32560.000000	32560.000000
<b>mean</b>	38.581634	1.897818e+05	10.080590	1077.615172	87.306511	40.437469
<b>std</b>	13.640642	1.055498e+05	2.572709	7385.402999	402.966116	12.347618
<b>min</b>	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
<b>25%</b>	28.000000	1.178315e+05	9.000000	0.000000	0.000000	40.000000
<b>50%</b>	37.000000	1.783630e+05	10.000000	0.000000	0.000000	40.000000
<b>75%</b>	48.000000	2.370545e+05	12.000000	0.000000	0.000000	45.000000
<b>max</b>	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000

```
In [267]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32560 entries, 0 to 32559
Data columns (total 15 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   age                   32560 non-null  int64
 1   workclass              32560 non-null  object
 2   fnlwgt                 32560 non-null  int64
 3   education              32560 non-null  object
 4   education-num          32560 non-null  int64
 5   marital-status         32560 non-null  object
 6   occupation             32560 non-null  object
 7   relationship           32560 non-null  object
 8   race                   32560 non-null  object
 9   sex                    32560 non-null  object
10   capital-gain           32560 non-null  int64
11   capital-loss           32560 non-null  int64
12   hours-per-week         32560 non-null  int64
13   native-country         32560 non-null  object
14   income                 32560 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
In [268]: data.isnull().sum()
```

```
Out[268]: age                0
workclass              0
fnlwgt                 0
education              0
education-num          0
marital-status         0
occupation             0
relationship           0
race                   0
sex                    0
capital-gain           0
capital-loss           0
hours-per-week         0
native-country         0
income                 0
dtype: int64
```

**there is no null values**

**splitting into numerical and categorical**

```
In [269]: numeric=data.select_dtypes(include=np.number).columns.tolist()
```

```
In [270]: numeric
```

```
Out[270]: ['age',
'fnlwgt',
'education-num',
'capital-gain',
'capital-loss',
'hours-per-week']
```

```
In [271]: category=data.select_dtypes(exclude=np.number).columns.tolist()
```

```
In [272]: category
```

```
Out[272]: ['workclass',  
           'education',  
           'marital-status',  
           'occupation',  
           'relationship',  
           'race',  
           'sex',  
           'native-country',  
           'income']
```

```
In [273]: num=data[numeric]  
          cat=data[category]
```

```
In [274]: num.head()
```

```
Out[274]:
```

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week
0	50	83311	13	0	0	13
1	38	215646	9	0	0	40
2	53	234721	7	0	0	40
3	28	338409	13	0	0	40
4	37	284582	14	0	0	40

```
In [275]: cat.head()
```

```
Out[275]:
```

	workclass	education	marital-status	occupation	relationship	race	sex	native-country	income
0	Self-emp-not-inc	Bachelors	Married-civ-spouse	Exec-managerial	Husband	White	Male	United-States	<=50K
1	Private	HS-grad	Divorced	Handlers-cleaners	Not-in-family	White	Male	United-States	<=50K
2	Private	11th	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	United-States	<=50K
3	Private	Bachelors	Married-civ-spouse	Prof-specialty	Wife	Black	Female	Cuba	<=50K
4	Private	Masters	Married-civ-spouse	Exec-managerial	Wife	White	Female	United-States	<=50K

## EDA

```
In [276]: import matplotlib.pyplot as plt
```

```

In [277]: # Function to perform univariate analysis of categorical columns
def plot_categorical_columns(dataframe):
    categorical_columns = dataframe.select_dtypes(include=['object']).columns

    for i in range(0, len(categorical_columns), 2):
        if len(categorical_columns) > i+1:

            plt.figure(figsize=(10,4))
            plt.subplot(121)
            dataframe[categorical_columns[i]].value_counts(normalize=True).plot(kind='bar')
            plt.title(categorical_columns[i])
            plt.subplot(122)
            dataframe[categorical_columns[i+1]].value_counts(normalize=True).plot(kind='bar')
            plt.title(categorical_columns[i+1])
            plt.tight_layout()
            plt.show()


        else:
            dataframe[categorical_columns[i]].value_counts(normalize=True).plot(kind='bar')
            plt.title(categorical_columns[i])

```

```

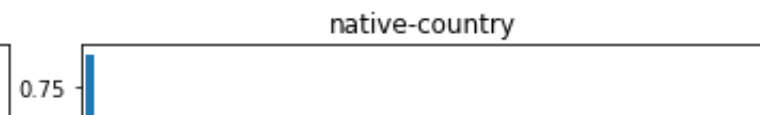




In [278]: plot_categorical_columns(cat)

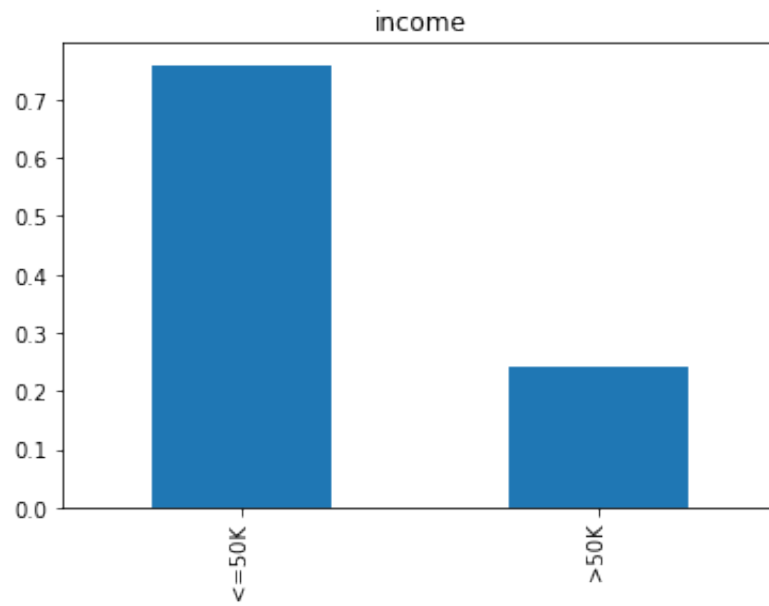

```





Se





```
In [279]: import seaborn as sns
```

```

In [280]: # Function to plot histograms
def plot_continuous_columns(dataframe):
    numeric_columns = dataframe.select_dtypes(include=['number']).columns.tolist()
    dataframe = dataframe[numeric_columns]

    for i in range(0, len(numeric_columns), 2):
        if len(numeric_columns) > i+1:
            plt.figure(figsize=(10,4))
            plt.subplot(121)
            sns.distplot(dataframe[numeric_columns[i]], kde=False)
            plt.subplot(122)
            sns.distplot(dataframe[numeric_columns[i+1]], kde=False)
            plt.tight_layout()
            plt.show()

        else:
            sns.distplot(dataframe[numeric_columns[i]], kde=False)

# Function to plot boxplots
def plot_box_plots(dataframe):
    numeric_columns = dataframe.select_dtypes(include=['number']).columns.tolist()
    dataframe = dataframe[numeric_columns]

    for i in range(0, len(numeric_columns), 2):
        if len(numeric_columns) > i+1:
            plt.figure(figsize=(10,4))
            plt.subplot(121)
            sns.boxplot(dataframe[numeric_columns[i]])
            plt.subplot(122)
            sns.boxplot(dataframe[numeric_columns[i+1]])
            plt.tight_layout()
            plt.show()





        else:
            sns.boxplot(dataframe[numeric_columns[i]])


```

```


In [281]: plot_continuous_columns(num)

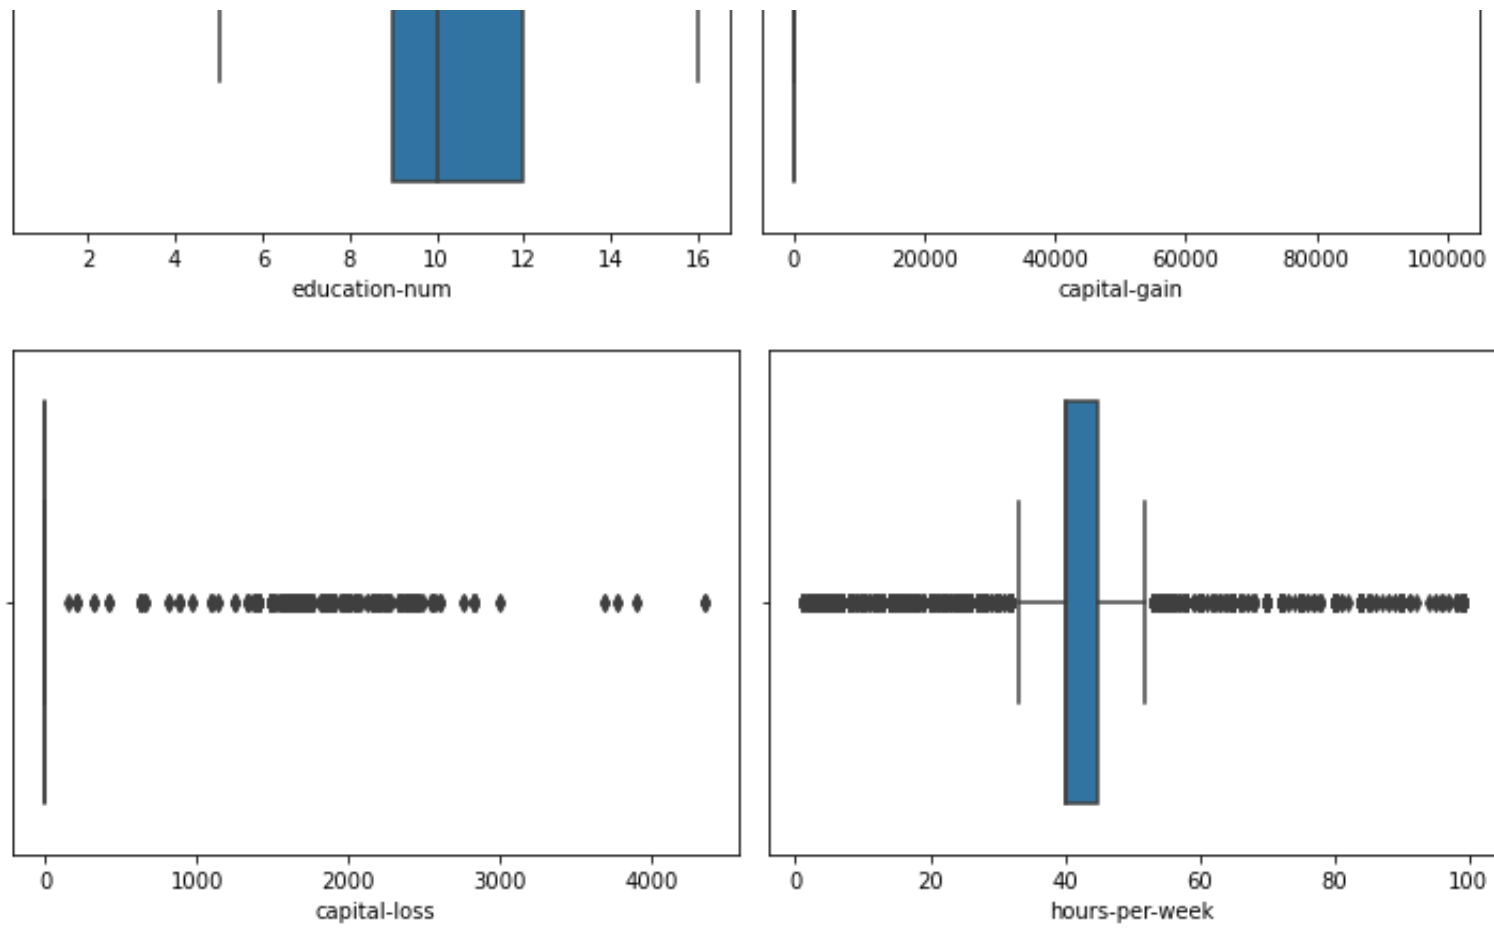
```





In [282]: `plot_box_plots(num)`





**treating the outliers**

```
In [283]: from scipy.stats.mstats import winsorize
# Function to treat outliers
def treat_outliers(dataframe):
    cols = list(dataframe)
    for col in cols:
        if col in dataframe.select_dtypes(include=np.number).columns:
            dataframe[col] = winsorize(dataframe[col], limits=[0.05, 0.1], inclusive=(True, True))

    return dataframe
```

```
In [284]: num=treat_outliers(num)
```

```
In [285]: data.income.unique()[0]
```

```
Out[285]: ' <=50K'
```

```
In [286]: data['income'] = data['income'].map({data.income.unique()[0]:0,data.income.unique()[1]:1})
```

```
In [287]: data.tail()
```

Out[287]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country
32555	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White	Female	0	0	38	United-States
32556	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	0	0	40	United-States
32557	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried	White	Female	0	0	40	United-States
32558	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	White	Male	0	0	20	United-States
32559	52	Self-emp-inc	287927	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White	Female	15024	0	40	United-States

```
In [298]: pd.crosstab(data['sex'],data['income'])/data.shape[0]*100
```

Out[298]:

income	0	1
sex		
Female	29.459459	3.621007
Male	46.458845	20.460688

clearly shows that male are highly employed and female employment rate is low

splitting into x and y




```
In [219]: x=data.drop(['income'],axis=1)
          y=data['income']
```

```
In [222]: data['income'].value_counts()
```

```
Out[222]: 0    24719
          1     7841
          Name: income, dtype: int64
```

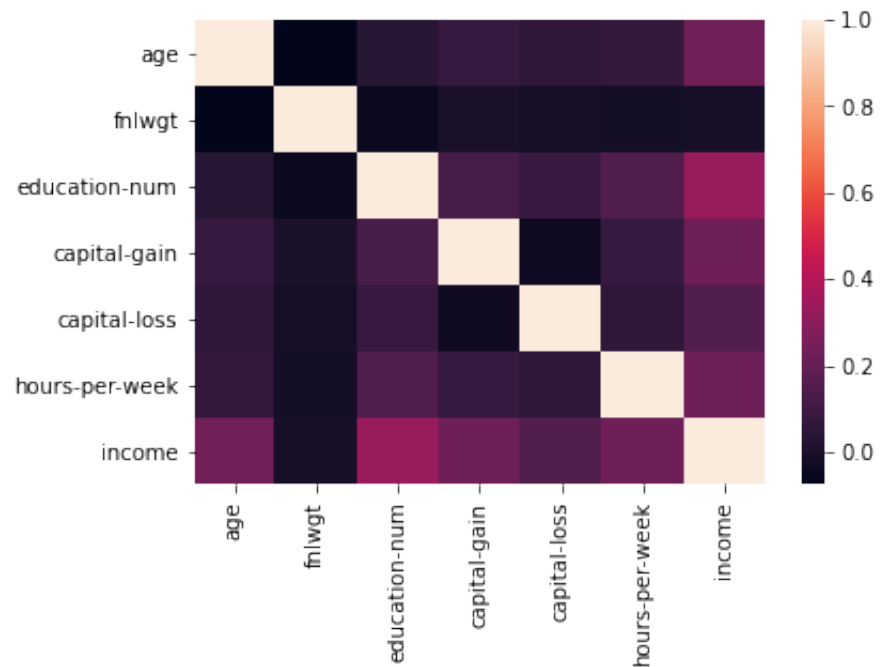
```
In [224]: data['income'].value_counts().plot.pie()
```

```
Out[224]: <AxesSubplot:ylabel='income'>
```



```
In [225]: #pca,linear assuption : no multi correality  
# from heatmap  
sns.heatmap(data.corr())
```

Out[225]: <AxesSubplot:>



## converting categorical into numerical data

```
In [239]: x=pd.get_dummies(x)
```

```
In [240]: x.head()
```

```
Out[240]:
```

workclass_ ?	workclass_ Federal- gov	workclass_ Local-gov	workclass_ Never- worked	...	native- country_ Portugal	native- country_ Puerto- Rico	native- country_ Scotland	native- country_ South	native- country_ Taiwan	native- country_ Thailand	native-country_ Trinidad&Tobago	native- country_ United- States	nat coun Vietn
0	0	0	0	...	0	0	0	0	0	0	0	1	
0	0	0	0	...	0	0	0	0	0	0	0	1	
0	0	0	0	...	0	0	0	0	0	0	0	1	
0	0	0	0	...	0	0	0	0	0	0	0	0	
0	0	0	0	...	0	0	0	0	0	0	0	1	

## Logistic Regression

```
In [241]: from sklearn.model_selection import train_test_split
```

```
In [242]: X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2)
```

```
In [243]: from sklearn.linear_model import LogisticRegression
```

```
In [244]: model=LogisticRegression()
```

```
In [245]: model.fit(X_train,y_train)
```

```
Out[245]: LogisticRegression()
```

```
In [246]: pred=model.predict(X_test)
```

```
In [255]: from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
```

```
In [256]: confusion_matrix(pred, y_test)
```

```
Out[256]: array([[4759, 1157],  
                [ 159,  437]])
```

```
In [257]: print(classification_report(pred, y_test))
```

	precision	recall	f1-score	support
0	0.97	0.80	0.88	5916
1	0.27	0.73	0.40	596
accuracy			0.80	6512
macro avg	0.62	0.77	0.64	6512
weighted avg	0.90	0.80	0.83	6512

```
In [258]: accuracy_score(pred, y_test)
```

```
Out[258]: 0.797911547911548
```

## NSE

```
In [91]: import pandas as pd  
import numpy as np
```

```
In [92]: import matplotlib.pyplot as plt
import pandas as pd

import datetime as dt
import numpy as np
import os
from sklearn.preprocessing import MinMaxScaler
```

```
In [93]: dt=pd.read_csv('/Users/persie/Downloads/banknifty.csv')
```

```
In [94]: dt.head()
```

Out[94]:

	index	date	time	open	high	low	close
0	BANKNIFTY	20121203	09:16	12125.70	12161.70	12125.70	12160.95
1	BANKNIFTY	20121203	09:17	12161.75	12164.80	12130.40	12130.40
2	BANKNIFTY	20121203	09:18	12126.85	12156.10	12126.85	12156.10
3	BANKNIFTY	20121203	09:19	12157.25	12164.75	12151.60	12164.20
4	BANKNIFTY	20121203	09:20	12162.80	12162.80	12148.20	12151.15

```
In [95]: dt.describe()
```

```
Out[95]:
```

	date	open	high	low	close
<b>count</b>	3.675750e+05	367575.000000	367575.000000	367575.000000	367575.000000
<b>mean</b>	2.014401e+07	15078.023296	15082.498465	15073.480983	15077.993028
<b>std</b>	1.169302e+04	3184.438089	3185.213591	3183.628315	3184.411825
<b>min</b>	2.012110e+07	1405.050000	1407.050000	1404.600000	1405.200000
<b>25%</b>	2.013103e+07	12092.200000	12095.000000	12089.150000	12092.175000
<b>50%</b>	2.014110e+07	15526.100000	15531.200000	15521.400000	15525.950000
<b>75%</b>	2.015103e+07	17956.050000	17960.550000	17951.100000	17955.800000
<b>max</b>	2.016093e+07	20903.950000	20907.550000	20899.250000	20907.550000

```
In [96]: dt.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 367575 entries, 0 to 367574
Data columns (total 7 columns):
#   Column  Non-Null Count  Dtype
---  -
0   index   367575 non-null   object
1   date    367575 non-null   int64
2   time    367575 non-null   object
3   open    367575 non-null   float64
4   high    367575 non-null   float64
5   low     367575 non-null   float64
6   close   367575 non-null   float64
dtypes: float64(4), int64(1), object(2)
memory usage: 19.6+ MB
```

```
In [97]: dt.isnull().sum()
```

```
Out[97]: index      0  
date        0  
time        0  
open        0  
high        0  
low         0  
close       0  
dtype: int64
```

In [98]: ##### PLOTTING GRAPH #####

```
plt.figure(figsize = (18,9))
plt.plot(range(dt.shape[0]),(dt['low']+dt['high'])/2.0)
# plt.plot(range(df.shape[0]),(df['volume']))
plt.xticks(range(0,dt.shape[0],10000),dt['date'].loc[:,10000],rotation=45)
plt.xlabel('Date',fontsize=18)
plt.ylabel('Mid Price',fontsize=18)
plt.show()
```





### Mid price range for every day

```
In [99]: dt=dt.drop(['index','date','time'],axis=1)
```

```
In [100]: y=dt["close"]
```

```
In [101]: y.head()
```

```
Out[101]: 0    12160.95  
          1    12130.40  
          2    12156.10  
          3    12164.20  
          4    12151.15  
          Name: close, dtype: float64
```

```
In [102]: x=dt.drop(['close'],axis=1)
```

```
In [103]: x.head()
```

```
Out[103]:
```

	open	high	low
0	12125.70	12161.70	12125.70
1	12161.75	12164.80	12130.40
2	12126.85	12156.10	12126.85
3	12157.25	12164.75	12151.60
4	12162.80	12162.80	12148.20

## Modelling

```
In [104]: from sklearn.model_selection import train_test_split
```

```
In [105]: X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2)
```

```
In [106]: print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
(294060, 3) (73515, 3) (294060,) (73515,)
```

```
In [107]: from sklearn.linear_model import LinearRegression
```

```
In [108]: model = LinearRegression()
model.fit(X_train, y_train)
```

```
Out[108]: LinearRegression()
```

```
In [109]: pred = model.predict(X_test)
```

```
In [110]: from sklearn.metrics import r2_score
r2_score(y_test, pred)
```

```
Out[110]: 0.9999990415317453
```

**accuracy = 99%**